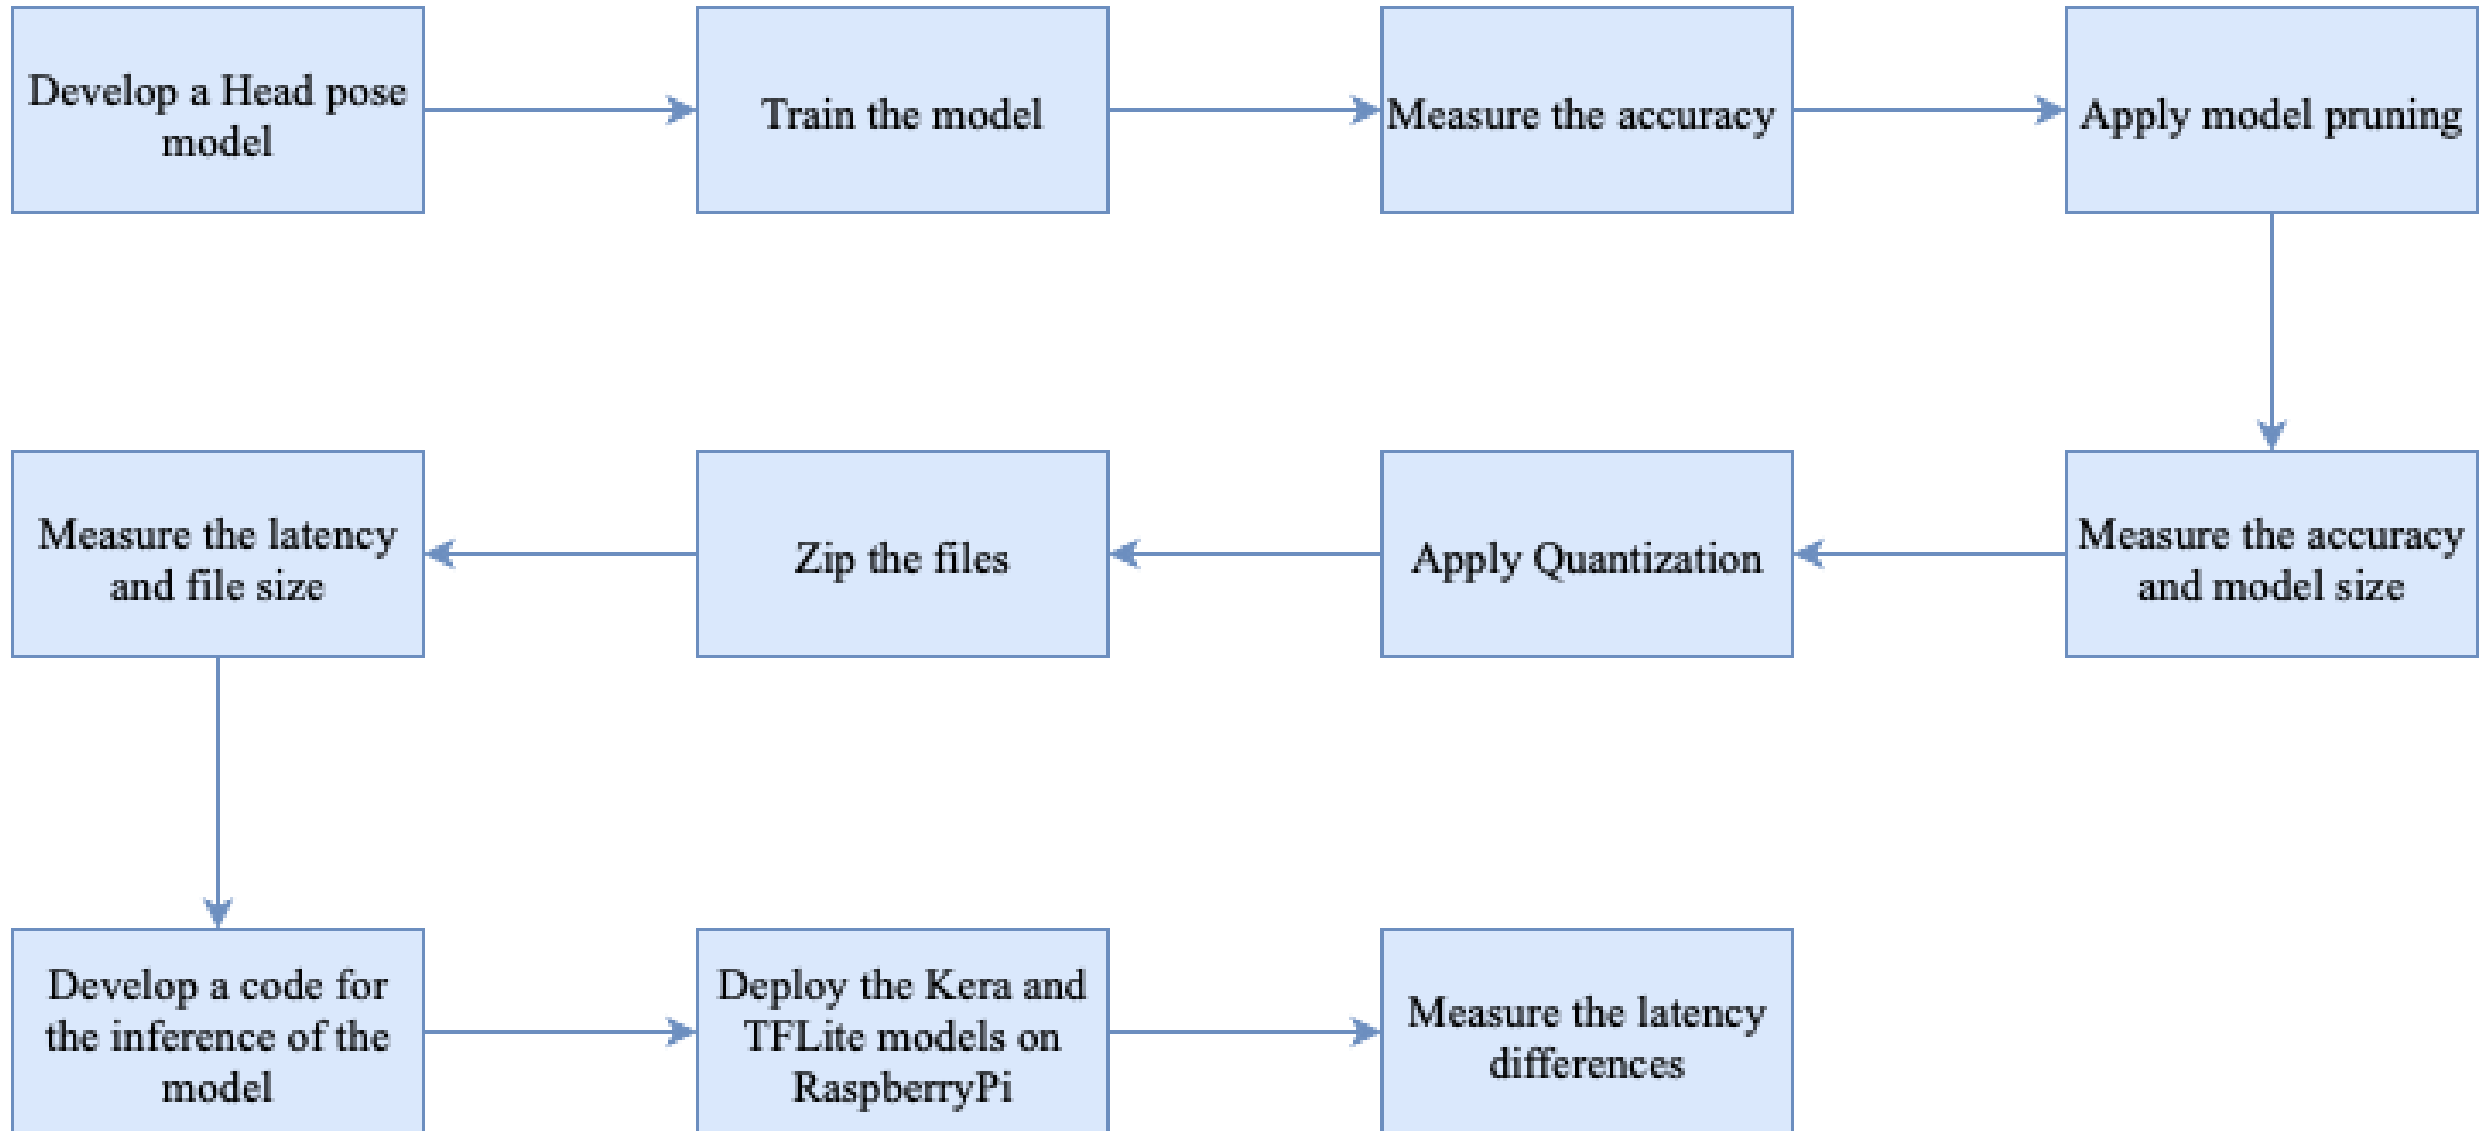


# TinyML: A Compact Revolution in Engineering AI

Session 3.1 : Hands-on session on Model Pruning and Quantization for head pose estimation using Tensorflow

Asiri Gawesha Lindamulage  
Open University of Sri Lanka

# End-to-end model deployment pipeline



# What is model pruning?


- Model pruning removes unnecessary weights from the neural network to,

- Improve the inference time
- Reduce File size


- Weights almost zero below a given threshold  $\rightarrow 0$

- Model pruning has two types,

- Dynamic pruning
- Static pruning



Before pruning



After pruning

$$s_t = s_f$$

where

$s_t$  = Current sparsity

$s_f$  = final sparsity

*Static sparsity equation*

$$s_t = s_f + (s_i - s_f) \left( 1 - \frac{t - t_0}{(t_f - t_0) \Delta t} \right)^3,$$

where

$s_t$  = Current sparsity value

$s_f$  = final sparsity

$t_f$  = Ending training step

$s_i$  = Initial sparsity

$t_0$  = Starting training step

$\Delta t$  = Pruning frequency

*Dynamic sparsity equation*

# Why Head Pose Estimation

Used in many applications


- AR/VR headsets and wearables
- portable assistive technologies
- Driver Monitoring Systems

Down Stream Tasks

- Gaze Estimation
- Human-Computer Interaction
- Action Recognition
- Regression method
- The Mean Absolute Error is calculated by averaging the errors of the three angles.

$$\text{MAE} = \frac{1}{3} (|y - \hat{y}| + |p - \hat{p}| + |r - \hat{r}|)$$

- Benchmarking Datasets- BIWI, AFLW2000, 300WLP



Aspect	Arduino Nano 33 BLE Sense	Raspberry Pi 4/5	Normal PC (Laptop/Desktop)
Processor	ARM Cortex-M4 @ 64 MHz	ARM Cortex-A72 (quad-core @ ~1.5,~2 GHz)	Intel/AMD x86 (multi-core, ~3+ GHz)
RAM	256 KB SRAM	1,~8 GB LPDDR4	4,~64 GB DDR4/DDR5
Flash/Storage	1 MB Flash	8,~128 GB microSD or SSD	256 GB ,~2 TB SSD/HDD
ML Capability	TinyML models only	Quantized and small full ML models	Full-scale ML/DL (CNNs, Transformers)
Model Format	TFLite Micro (.tflite)	TFLite / ONNX / PyTorch (lightweight)	TensorFlow, PyTorch, ONNX, JAX
Real-Time Support	Excellent (<1ms latency)	Moderate (10,~100ms latency)	Variable (not real-time by design)
Power Consumption	~20 mW	~2.5,~10 W	~30,~300+ W
Battery Operable	Yes (coin cell or LiPo)	Yes (power bank)	Limited (laptop), No (desktop)



# Some experimental results

- Model Used- EfficientNet(B0) based Headpose Estimation model [1]
- Trained on BIWI Kinect Head Pose [13], 300W-LP [14] and tested on AFLW 2000 [14]
- 8.63 Mean absolute error (MAE) with a file size of 110MB
- Analyse the parameter selection for model pruning
  - Static pruning
    - Static sparsity - 50%, 75% and 87.5%
    - Starting step (epoch) - 0, 20, 40, 60, 80
    - Ending step (epoch) – 20, 40, 60, 80
  - Dynamic pruning
    - Final sparsity – 50%, 75%, and 87.5%
    - Starting step (epoch) - 0, 20, 40, 60, 80
    - Ending step (epoch) – 20, 40, 60, 80
- Combine pruning with post optimizers
  - Experimental post pruning optimizer
  - Dynamic post quantization




Figure 1. Model Architecture (source-[12])



# MAE and model size variation with sparsity percentage




Figure 1. Dynamic Pruning







Figure 2. Static Pruning

- Sparsity percentage  $\uparrow$   $\rightarrow$  MAE  $\uparrow$  - Accuracy  $\downarrow$  - File size  $\downarrow$
- Post pruning optimizer reduces file size without altering the model accuracy
- Dynamic quantization can reduce file size significantly with a slightly reduced model accuracy



# Starting step and ending step



- Best starting point = 0
- Best ending point = 60 (3/4 of the total training steps)
- Static and dynamic pruning behaviour is identical
- Post optimizers:
- Post pruning optimizer reduces model size around 3 MB without altering model accuracy
- Dynamic post quantizer reduces the file size upto 15MB while sacrificing model accuracy below 0.1 (MAE)



# Summary of the results

BEST MODELS OF DYNAMIC AND CONSTANT PRUNING

Model	Training parameters				Pruned model		Post pruned model		Post quantized model	
	Initial Sparsity	Final Sparsity	Starting step	End step	MAE	Size (MB)	MAE	Size (MB)	MAE	Size (MB)
Dynamic(Best Accuracy)	0.00	0.50	0.00	60.00	8.54	22.08	8.54	21.82	8.58	6.24
Dynamic(Best model size)	0.00	0.88	40.00	80.00	12.31	9.22	11.98	8.58	12.17	2.54
Constant((Best Accuracy)	0.00	0.50	60.00	80.00	8.57	35.23	8.57	35.24	8.59	7.85
Constant(Best model size)	0.00	0.88	60.00	80.00	11	9.21	12.39	8.41	12.51	2.59

- Starting step and ending step – best combination 0 to 60.
- Final sparsity – final sparsity higher -> lower the model accuracy
  - If priority -> model size - 87.5%
  - If priority-> model accuracy - 50%
- File size can be further reduced by combining with dynamic post- quantization sacrificing model accuracy

Source: Comparative Study of Parameter Selection for Enhanced Edge Inference for a Multi-Output Regression model for Head Pose Estimation (Tencon 2022)





# TinyML: A Compact Revolution in Engineering AI

Session 3.2 : Hands-on session on deploying a simple CNN on Arduino BLE 33

Asiri Gawesha Lindamulage

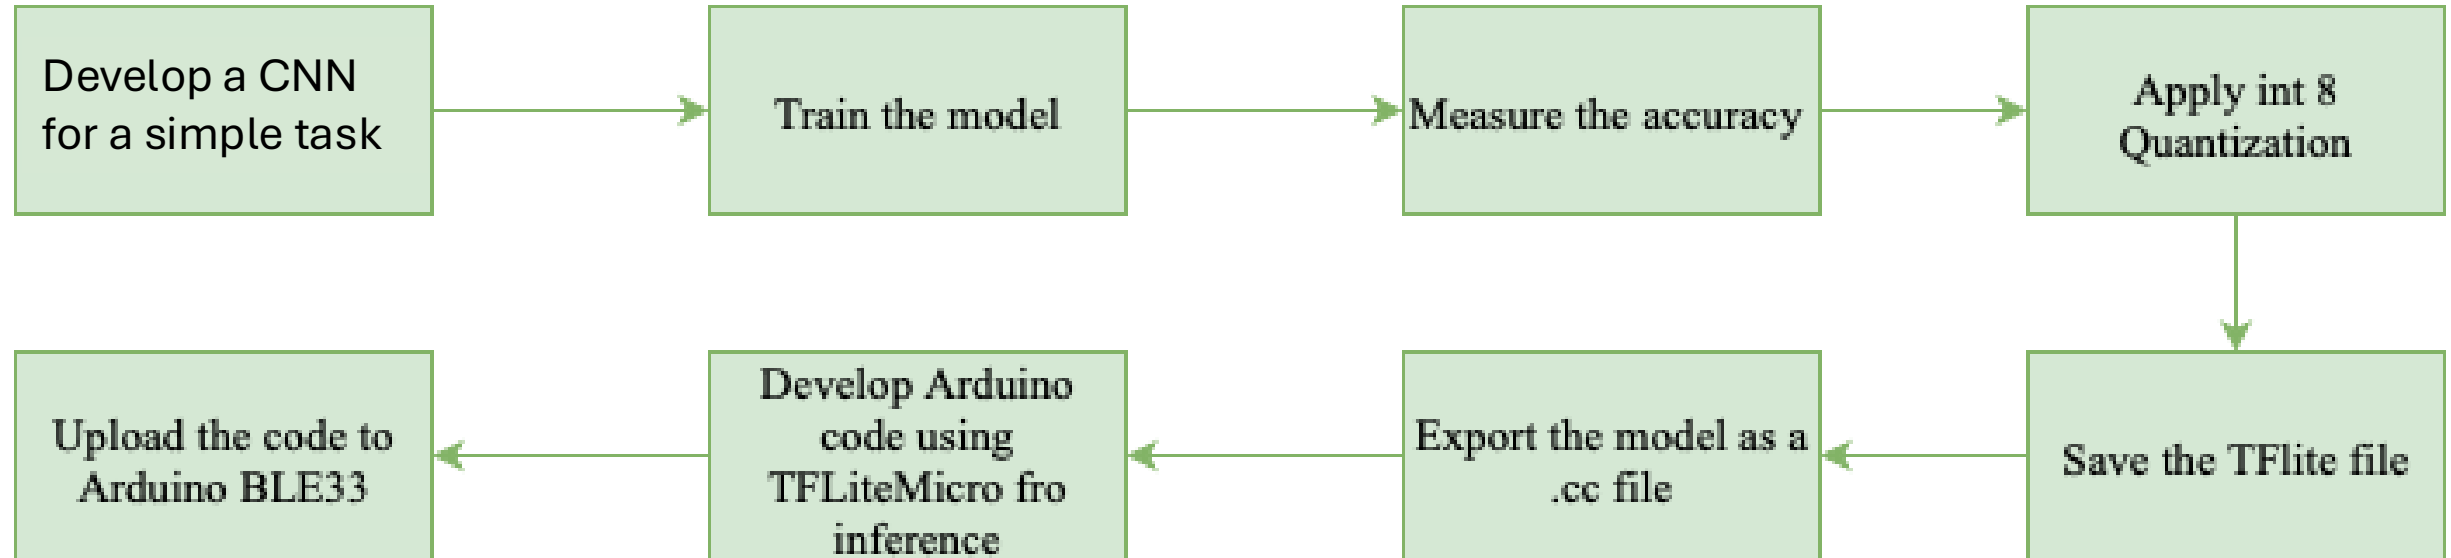
Open University of Sri Lanka

aglin@ou.ac.lk



THE  
**OPEN UNIVERSITY**  
OF SRI LANKA

# Arduino BLE 33 Deployment Cycle



# Some very recent structured model pruning techniques

Some of the key turning points in developement of model pruning techniques

Theories	Authors
To prune, or not to prune: exploring the efficacy of pruning for model compression – Unstructured pruning method.	Michael Zhu, Suyog Gupta
THE LOTTERY TICKET HYPOTHESIS (2019) – finds a smaller architecture hidden within the original model	Jonathan Frankle, Michael Carbin
DepGraph: Towards Any Structural Pruning (2023) - based on the connections between layers and its been applied to GNNs and CNNS	Gongfan Fang, Xinyin Ma, Mingli Song, Michael Bi Mi, Xinchao Wang
Optimized Transformer Models: l'BERT with CNN-like Pruning and Quantization(2024)	Muhammad Hamis Haider, Stephany Valarezo-Plaza, Sayed Muhsin
ARPruning: An automatic channel pruning based on attention map ranking(2024)	Tongtong Yuan , Zulin Li , Bo Liu , Yinan Tang , Yujia Liu



# Experimental possibilities

- Combine model pruning with quantization.
- Try structured model pruning methods available on PyTorch
- Develop structured model pruning techniques

